

---

# Amazon Product Recommendations Using Collaborative Filtering and Natural Language Processing

---

**Nikhil Aourpally**  
Nikhil.Aourpally@asu.edu

**Anthony Helmstetter**  
anthony.helmstetter@asu.edu

**Prajwal Paudyal**  
ppaudyal@asu.edu

**Anirudh Som**  
Anirudh.Som@asu.edu

## Abstract

In this project we will implement a Recommender System to provide suggestions based on Amazon user reviews across Amazon Instant Video, Music, and Book product categories. The Recommender System will be implemented using collaborative filtering based on product ratings. To increase recommendation accuracy, we will also perform natural language processing on user written reviews to better inform the recommender system of user preferences. We will assess the accuracy of the Recommender System with and without additional textual information contained in user reviews, and evaluate the sensitivity of our Recommender System by varying the number of reviews by user, and number of reviews per product.

## 1 Introduction

Modern consumers are inundated with a plethora of choices. Electronic retailers and content providers offer a huge selection of products, with unprecedented opportunities to meet users having a variety of special needs and tastes. Matching consumers with most appropriate products is key to enhancing user satisfaction and loyalty. For this reason, many online retailers have become interested in recommender systems, which analyze patterns of user interest in products. This allows them to provide personalized recommendations that suit each and every user's taste.

Personalized recommendation can add another dimension to a user's online shopping experience. Giants like Amazon.com, Netflix and other e-commerce leaders have made recommender systems a salient part of their websites. Recommender systems can be treated as a form of information filter. In general, a recommender system needs to complete two tasks [3]- First, we use a recommendation algorithm and the user's history logs to predict item ratings[4], which is not marked by the user before. Second [1], we generate a list of recommendation items for the user on the basis of the first task mentioned before, and also on a number of other factors like diversity, novelty, credits and so on. For every recommendation system, the first period is key and the core technology is the recommendation algorithm.

## 1.1 Collaborative Filtering

Recommender systems are based on one of two strategies [5], namely - content filtering and collaborative filtering (CF). The content filtering approach creates a profile for each user or product, in order to characterize its nature. These profiles allow programs to associate users with matching products. However, this strategy requires gathering external information that might not be available or easy to collect. On the other hand, collaborative filtering relies on past user behavior i.e. previous transactions or product ratings without requiring the creation of explicit profiles. Collaborative filtering analyzes relationships between users and interdependencies between products, to identify new user-item associations. The goal of collaborative filtering is to suggest the new items to users by predicting a score for an item (usually 0/1 or 1 – 5, etc) for a particular user based on his consume log of some old items as well as the log from other users with similar tastes.

Memory-based CF has two models - one is user-based and the other is item/product-based. The basic idea of user-based CF can be explained using the following example - suppose user A likes an item a, user B likes items a, b, c and user C likes items a and c. So user A, B and C are similar and item c is recommended to user A. Item-based CF calculates the similarity between items with all the history of the user's preference data, and thus the user is recommended with similar items. Model-based CF are developed using data mining, machine learning algorithms to find patterns based on training data which are used to make predictions for real data. Most of the recommender system algorithms have been derived from collaborative filtering. The biggest advantage of CF is that it is domain free. However, it suffers from cold start problems, due to its inability to address the system's new products and users. In this aspect, content filtering is superior.

## 1.2 Matrix Factorization

The primary areas of collaborative filtering [5] are the neighborhood methods and latent factor models. Neighborhood methods are centered on computing the relationships between items or, alternatively, between users. The item oriented approach evaluates a user's preference for an oriented item based on ratings of neighborhood items by the same user. A product's neighborhood are other products that tend to get similar ratings when rated by the same user. Latent factor models are an alternative approach that tried to explain the ratings by characterizing both items and users on, say 20 to 100 factors inferred from the ratings patterns. Most of the successful realizations of latent factor models are based on matrix factorization (MF).

The MF technique is gradually applied to personalized recommendation algorithm. It does a global optimization on object, so MF performs better than neighborhood methods in Root Mean Square Error (RMSE). In the basic form, MF characterizes both items and users by vectors of factors inferred from item rating patterns. High correspondence between item and user factors leads to a recommendation. One strength of MF is that it allows incorporation of additional information. When explicit feedback is not available, recommender systems can infer user preferences using implicit feedback, which indirectly reflects opinion by observing user behavior including purchase history, browsing history, search patterns, or even mouse movements. Implicit feedback usually denotes the presence or absence of an event, so it is typically represented by a densely filled matrix. However, ideally the MF approach will always prefer explicit feedback than implicit.

MF models map both users and items to a joint latent factor space of dimensionality  $f$ , such that user-item interactions are modeled as inner products in that space. Accordingly, each item  $i$  is associated with a vector  $q_i \in \mathbb{R}^f$ , and each user  $u$  is associated with a vector  $p_u \in \mathbb{R}^f$ . For a given item  $i$ , the elements of  $q_i$  measure the extent to which the item possesses those factors, positive or negative. For a given user  $u$ , the elements of  $p_u$  measure the extent of interest the user has in items that are high on the corresponding factors, again, positive or negative. The resulting dot product,  $q_i^T \cdot p_u$ , captures the interaction between user  $u$  and item  $i$  - the user's overall interest in the item's characteristics. This approximates user  $u$ 's rating of item  $i$ , which is denoted by  $r_{ui}$ , leading to the estimate

## EQUATION 1

To learn the factor vectors ( $p_u$  and  $q_i$ ), the system minimizes the regularized squared error on the set of known ratings:

## EQUATION 2

Here,  $\kappa$  is the set of the  $(u, i)$  pairs for which  $r_{ui}$  is known (training set). The system learns the model by fitting the previously observed ratings. However, the goal is to generalize those previous ratings in a way that predicts future, unknown ratings. Thus, the system should avoid overfitting the observed data by regularizing the learned parameters, whose magnitudes are penalized. The constant  $\lambda$  controls the extent of regularization and is usually determined by cross-validation.

### 1.3 Alternating Least Squares

There are two approaches [5] to minimize equation2 stochastic gradient descent and alternating least squares (ALS). Both  $q_i$  and  $p_u$  are unknowns, and equation2 is not convex. By fixing one of the unknowns, the optimization problem becomes quadratic and can be solved optimally. Thus the ALS techniques rotate between fixing the  $q_i$ 's and fixing the  $p_u$ 's. When all  $p_u$ 's are fixed, the system recomputes the  $q_i$ 's by solving a least-squares problem, and next when all the  $q_i$ 's are fixed, the system recomputes the  $p_u$ 's by solving the least squares problem. This ensures that each step decreases equation2 until convergence.

In general, stochastic gradient descent is easier and much faster than ALS. However, ALS is more favorable than stochastic gradient descent in at least two cases ? First, when the system computes  $q_i$  and  $p_u$  independently of the other user factors respectively, the system can utilize massive parallelization of the algorithm. Next, when the system is centered on implicit data, because the training set cannot be considered sparse, and looping over each single training case as gradient descent does, would not be practical.

### 1.4 Natural language processing

Natural language processing (NLP) is a field of computer science, artificial intelligence, and linguistics concerned with the interactions between computers and human (natural) languages. Modern NLP algorithms are based on machine learning, especially statistical machine learning. Prior implementations of language-processing tasks typically involve the direct hand coding of large sets of rules. Some of the tasks in NLP have direct real-world applications, while other more commonly serve as subtasks that are used to aid in solving larger tasks. Goal of NLP evaluation is to measure one or more qualities of an algorithm or a system, in order to determine whether (or to what extent) the system answers the goals of its designers, or meets the needs of its users.

## 2 Problem Description

The rapid growth of internet and e-commerce has brought the problem of information overload. Excessive information makes it difficult for a user to pick up what they really want and reduces the efficiency of information usage. Several search engines or e-commerce online platforms use personalized recommender systems in searching results. Personalized recommender systems not only performs better than the traditional ones but also adds another dimension to the user's experience, by analyzing patterns of the user's interest in products to suit their taste. Using Amazon product reviews, we implemented a recommender system to suggest new movies to Amazon users. The suggested movies are determined by highly rated products as rated by similar users. The similarity of users will be determined by the degree to which the users have rated identical products similarly. Our recommender system falls under the collaborative filtering category, in which we filter based on the similarity of users and not by the similarity of products.

To enhance the performance of our recommender system, we will also incorporate additional information derived from text reviews written by users using Natural Language Processing, to develop a more accurate metric of user similarity. We will evaluate the performance of our recommender system both with and without the incorporation of additional textual information using NLP and compare the results.

### 3 Methodology

#### 3.1 Dataset

The Dataset consists reviews for movies and television shows available on the Amazon Instant Video service. The data is publically available at <https://snap.stanford.edu/data/web-Amazon.html>. The Amazon-Instant-Video is 252 MB and consists of 717,651 reviews across 312929 unique users and 22190 unique movies. A typical entry is shown below.

```
product/productId: B002BNZ2XE
product/title: Amazon.com: Moving Target (1988): Jason Bateman, John Glover,
Jack Wagner, Chynna Phillips: Amazon Instant Video
product/price: 0.00
review/userId: AZVY9Y3A0YU1I
review/profileName: Willard R. Stephen
review/helpfulness: 1/1
review/score: 4.0
review/time: 1277424000
review/summary: Moving Target
review/text: I found this to be an intrieging suspense thriller with a
minimum of violence and a happy ending. This is a rare combination in
recent releases. I found it well excepted and enjoyed by teen age
audiences.
```

The features of interest in the recommender system will be productId, userId, score, and review text. The score, or rating, of a review is an integer value between 1 and 5 inclusive. The review text consists of natural language text of variable length.

#### 3.2 Preliminary Neighborhood Methods

Our initial goal was to establish a baseline recommender system without incorporating a natural language processing component. To that end, we first explored neighborhood based collaborative filtering determined by user similarity. Using a smaller, more manageable subset of the the Amazon Instant Video dataset we computed the similarity between each pair of users. To calculate the similarity between users  $u$  and  $v$ , we first take the intersection of all movies rated by both  $u$  and  $v$  and create ratings vectors  $r_u$  and  $r_v$  whose elements consist of the explicit ratings by each user, for each movie in the shared intersection. Similarity between users was then calculated by computing the Pearson correlation coefficient between users.

After all pairwise similarity measures between users were calculated, recommendations for user  $u$  could be found by selecting the  $k$  most similar users. The predicted rating for a particular movie  $i$ , previously unrated by user  $u$  is then calculated by a weighted combination of the  $k$  most similar users' rating of movie  $i$ .

As the full Amazon Instant Video Set consisted of approximately 60,000 unique users, computing the similarity matrix would require over 2 billion function calls to compute the Pearson correlation coefficient. Procedures to parallelize these calcuations were infeasible, and we explored other options such as matrix mactorization methods.

### 3.3 Matrix Factorization Methods

Matrix factorization methods offer many advantages over neighborhood methods, including speed, parallelization, accuracy, and explanatory power. To build a recommender system based on matrix factorization, we transformed a flat list of movie ratings seen in the table on the left, below, into an  $m \times n$  matrix depicted below on the right.

UserId	MovieId	Rating
1	1	5
1	4	2
2	1	4
2	2	3
3	5	1
4	2	1
4	5	4
$\vdots$	$\vdots$	$\vdots$

	MovieId					
UserId	1	2	3	4	5	...
1	5	0	0	2	0	...
2	4	3	0	0	0	...
3	0	0	0	0	1	...
4	0	1	0	0	4	...
5	5	4	0	0	3	...
6	0	3	2	3	0	...
$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\vdots$	$\ddots$

The movie rating matrix,  $R \in \mathbb{R}^{m \times n}$  has rows corresponding to users, columns corresponding to movies, where elements  $r_{ui}$  of give the value on a 1-5 scale given by user  $u$  to movie  $i$ . This matrix is extremely incomplete, and provides the rationale for constructing a recommender system in the first place. All ratings for which we have no data are given a value of 0, resulting in an extremely sparse matrix.

The goal of a matrix factorization recommender is to discover the matrices  $P \in \mathbb{R}^{m \times k}$  and  $Q \in \mathbb{R}^{k \times n}$  such that  $R \approx PQ$ . That is, we wish to minimize the norm difference between the known rating  $r_{ui}$  and the predicted rating  $\hat{r}_{ui}$  where the predicted rating is given by  $\hat{r}_{ui} = q_i^T p_u$  with  $q_i$  and  $p_u$  the  $i$ th column and  $u$ th row of  $Q$  and  $P$  respectively. The parameter  $k$  determines the effective rank of the matrix factors with  $k \ll m, n$ . The value of  $k$  determines the dimension of the lower dimensional subspace of latent features both in user space and item space that best retains the interactions between users and items. The vector  $p_u \in \mathbb{R}^{1 \times k}$  measures the extent to which user  $u$  possesses the  $k$  latent user features, while the vector  $q_i \in \mathbb{R}^{k \times 1}$  measures the extent to which item  $i$  possesses the  $k$  latent item features. The interaction between these two vectors, and the predicted rating  $\hat{r}_{ui}$  is given by the inner product of the two feature vectors. The factorization is presented visually below:

$$\begin{array}{c} \text{rating} \\ \text{matrix} \end{array} \begin{bmatrix} r_{11} & r_{12} & \cdots & r_{1N} \\ r_{21} & r_{22} & \cdots & r_{2N} \\ \vdots & \vdots & \ddots & \vdots \\ r_{M1} & r_{M2} & \cdots & r_{MN} \end{bmatrix} = \begin{array}{c} \text{user} \\ \text{features} \end{array} \begin{bmatrix} p_{11} & p_{12} & \cdots & p_{1K} \\ p_{21} & p_{22} & \cdots & p_{2K} \\ \vdots & \vdots & \ddots & \vdots \\ \vdots & \vdots & \ddots & \vdots \\ p_{M1} & p_{M2} & \cdots & p_{MK} \end{bmatrix} \begin{array}{c} \text{item} \\ \text{features} \end{array} \begin{bmatrix} q_{11} & q_{12} & \cdots & \cdots & \cdots & q_{1N} \\ q_{21} & q_{22} & \cdots & \cdots & \cdots & q_{2N} \\ \vdots & \vdots & \ddots & \cdots & \cdots & \vdots \\ q_{M1} & q_{M2} & \cdots & \cdots & \cdots & q_{KN} \end{bmatrix}$$

We wish to minimize the squared difference between observed ratings and predicted ratings. Let  $\kappa$  be the set of ordered pairs  $(u, i)$  for which we have observed ratings between user  $u$  and item  $i$ . To avoid overfitting the

training data, we introduce a regularization parameter  $\lambda$  to balance the norm size of the factorization vectors  $p$  and  $q$  against the square of the ratings residuals.  $P$  and  $Q$  then solve regularized least squares formulation:

$$(Q, P) = \arg \min_{q, p} \sum_{i, u \in \kappa} (r_{ui} - q_i^T p_u)^2 + \lambda(\|q\|^2 + \|p\|^2) \quad (1)$$

As both  $p_u$  and  $q_i$  are unknown there is no closed form solution that minimizes the above functional. To compute the desired factorization we use an Alternating Least Squares approach. In the iterative ALS approach we fix  $p_u$  and minimize the objective function with respect to the  $q_i$ . Then we fix  $q_i$  and minimize the objective function with respect to  $p_u$ , alternating as we proceed. We are then minimizing the pair of equations, each with respect to one variable:

$$J(p_u) = (r_u - p_u Q) W_u (r_u - p_u Q)^T + \lambda p_u p_u^T$$

$$J(q_i) = (r_i - P q_i) W_i (r_i - P q_i)^T + \lambda q_i q_i^T$$

with the binary diagonal weighting matrices  $W_u$  and  $W_i$

$$w_{ui} = \begin{cases} 1 & : \text{ user } u \text{ rated movie } i \\ 0 & : \text{ else } \end{cases}$$

The original minimization problem is over the sum of  $(u, i) \in \kappa$ , but with the above weighting matrices, the alternating least squares solution can be found over all  $(u, i)$ . The minimization of the above equations has closed formed solutions

$$p_u = (Q^T W_u Q + \lambda I)^{-1} Q^T W_u r_u$$

$$q_i = (P^T W_i P + \lambda I)^{-1} P^T W_i r_i$$

To proceed with ALS, we assign  $p_u$  a random initial value, and used the closed form solution to the regularized least squares problem in one variable, to calculate  $q_i$ . With  $q_i$  fixed, we now compute the updated  $p_u$  and continue alternating between the two unknowns until convergence or a maximum number of iterations has been achieved.

### 3.4 Natural Language Processing

Most eCommerce retailers provide users an option to leave a review for products they have purchased. These reviews not only provide future users an overview of an item, but they also contain valuable information that the retailer can use. User reviews generally come with a text review. The data that we deal with comes with two sections of text reviews: the review summary and the text review (detailed). While the users provide a numerical review, which is what we base our main recommender system on, it is sufficiently clear that the review texts and summaries also contain some valuable information. We can leverage this information to give a finer grained score to each review.

Thinking intuitively, when a person looks at a movie to decide on whether or not to watch it, the recommender that is built on review scores gives some information but it misses out on the details that is contained in the text reviews. The goal behind using NLP in our system is so that we can somehow use this information in the review summary and text to fine-tune the review scores that go into the recommender system. For instance for two reviews that have the same initial rating of 4, one review summary and review text might be very positive, while the other one might be mediocre. We want a movie recommender system that is able to mine the information in the reviews and give a recommendation that is appreciative of this extra information.

While there are many recommendation systems that exist out there, a novelty of this recommendation system is that it finds a way to incorporate the sentiment information in the reviews that others don't take into

consideration. A thing to be noted however is that although we use NLP to fine-tune the recommender, the results are not readily testable since we will ultimately compare against the numeric reviews in the test set, thus we don't expect the NLP system to do better than the base system in that respect. However we do fine tune our NLP system and the combination algorithms such that the overall test error is minimized.

**Stanford NLP.** The first method that we tried was the sentiment Analyzer that is provided by Stanford NLP lab [2]. This sentiment analyzer utilizes a 'Semantic Treebank' that includes fine grained sentiment labels for 215,154 phrases in the parse trees of 11, 855 sentences. This technique claims to increase the accuracy of predicting-fine grained sentiment labels for all phrases to 80.7 percent. This technique also successfully detects negation in a sentence which is very important for sentiment analysis on user reviews. After compilation of the NLP files, we get a 'pipeline' which takes a single sentence at a time and gives back a score that is between 0 and 4 which is convenient because we could scale this and combine it directly with our base recommender system. We built a program that would first get a sentiment value for the Review Summary section, then it would parse the Review Text section into individual sentences, get a sentiment value for each of those sections and then combine them and normalize it. We gave a slightly higher score to the Review Summary section than the review score section based on the formula:

```
java -cp "*" -Xmx2g edu.stanford.nlp.pipeline.StanfordCoreNLP -annotators tok-
enize,ssplit,pos,lemma,ner,parse,dcoref -file input.txt
```

$$(1.2 * ReviewSummaryScore + ReviewScore)/2.2$$

Although this approach seems to give us good results, we had to abandon it ultimately because it was slow and it could not give us results in a timely manner. The system took about 1 second to give a score back for each sentence review. We were processing about 600,000 reviews, which roughly translates to 6 million sentences, which amounts to about 70 days of runtime. On top of that there was the overhead of combination and some isolated cases where the sentence spanned more than 60 words which exponentially increased the processing time.

**Python Blob.** The next approach we utilized for natural language processing was TextBlob [6]. It provides a simple API for diving into common natural language processing (NLP) tasks such as part-of-speech tagging, noun phrase extraction, sentiment analysis, classification, translation, and more. Then we used the sentiment analysis module of this to receive a sentiment score between -1 and 1. The module returned a tuple of the form Sentiment(polarity,subjectivity). The polarity score is a float within the range [-1.0, 1.0]. The subjectivity is a float within the range [0.0, 1.0] where 0.0 is very objective and 1.0 is very subjective. In did not utilize the objectivity score in our NLP processing primarily because we did not want to differentiate between a negative reviews that was very objective vs. one that was subjective. (Example. The movie was awful vs. I felt the movie was awful). However, after going through several movie reviews, we realized that a good portion of the reviews were plot descriptions and not actual sentiments on the movie. In this light, giving more weight to reviews with higher subjectivity might increase the effectiveness of the sentiment analyzer. Due to time constraints, we were not able to run a detailed test on this aspect and had to set this aside for future work.

As compared to the Stanford NLP processor, the textBlob algorithm was very fast. We used a similar method to parse the Review Summary vs. the Review Text and fed it to the textBlob. We also noticed that this NLP system was able to handle entire paragraphs at a time, which was one of the reasons it was much faster. This algorithm was able to compute the entire training set of close to 600,000 reviews in less than 29 minutes. The result of the analysis was a sentiment score of between -1 and 1. We used the following conversion method to convert this range to a range of 1 to 5 using the formula

$$Score(1 - 5) = 2 * score + 3$$

**Combination.** We noticed that the sentiment score that was returned was somewhat positively skewed and thus shifted the overall scores upwards instead of just providing a means to differentiate between reviews

with the same initial score, thus we performed a simple mean shift to the NLP scores (shifted the mean to the integer scores). This made a lot of intuitive sense and also helped us reduce the testing error after we utilized the NLP system.

–insert diagrams for the mean shift.

We then used this shifted NLP score and combined based on the following formula.

$$BR + \alpha((AdNLP - BaseNLP)/4)$$

The  $\alpha$  here is the limiting factor which essentially limits how big of an adjustment that the NLP is allowed to have. For a value of 1, the NLP portion will be able to adjust the base score by at most 1 numeric score i.e. a review of 2 with the highest possible review score will be adjusted to a new score of 3. We performed a lot of tests on what would be the optimum value of the  $\alpha$  and settled on a value of 0.5. Since the reviews numeric score and the sentiment score didn't differ by such a huge degree in most cases, which meant that most of the adjustments stayed well below 0.5.

$$(1.2 * ReviewSummaryScore + ReviewScore)/2.2$$

## 4 Results

We split the dataset into 6 folds and we run 5-fold cross validation on the dataset. We obtained optimal values for the parameters rank, Iterations and regularization parameter  $\lambda$ . We performed 5-fold cross validation with rank being [10,20,40,60,80], iteration vector being [5,10,15,20] and  $\lambda$  vector being [10.0,1.0,0.1,0.01,0.001]. We performed 5-fold cross validation for various combination of rank, iteration and regularization parameter  $\lambda$ . \*Figure\* shows the variation of Mean Squared Error(MSE)(testing and training) with  $\lambda$ . \*Figure\* shows the variation of MSR with rank and \*Figure\* shows the variation of MSE with iterations. We achieved best testing error of 0.962879 for a rank, iterations and regularization parameter of 80,20 and 0.15 respectively. \*\*Reasons to round the predictions\*\* .After rounding the predictions, we obtained a MSE of 1.02 for the base recommender.

## 5 Future work

For the future, we can use a better Algorithm to do the NLP on the review text. Most of the users in the Amazon Instant Video Dataset that we used describe the plot of the movie in the review text rather than their opinion about the movie. So we think that we will be able to get a better insight into the review if we use dataset like Electronics where we have a subjective opinion of the users.

## 6 Conclusions

### References

- [1] Chuunhong Zhang, Bin Wang, Qing Liao. Weight based knn recommender system., 2013.
- [2] R. B. Bird, E. A. Smith, and D. W. Bird. The hunting handicap: costly signaling in human foraging strategies. *Behavioral Ecology and Sociobiology*, 50:9–19, 2001.
- [3] A. Tuzhilin, G. Adomavicius. Recommender systems handbook., 2010.
- [4] M. Reinders, J. Wang, A. Vries. A user-item relevance model for log-based collaborative filtering.
- [5] Chris Kolinsky, Yehudi Koren, Robert Bell. Techniques for recommender systems. *ATT*, page 2, 2009.
- [6] Chris Kolinsky, Yehudi Koren, Robert Bell. Techniques for recommender systems. *BTS*, 1:2, 2009.